

**CIS 4130- Big Data Technologies**

**Naziha Malik**

**[naziha.malik@baruchmail.cuny.edu](mailto:naziha.malik@baruchmail.cuny.edu)**

## **Milestone 1- Introduction/ Proposal**

For the semester-long project, I am proposing to use a set that describes the Uber and Lyft data in NYC from the years 2019- 2022. This data set includes very granular information about each ride, such as the TLC license plate, the date, the pick-up and drop off-times, total miles, total time of trip, and much more. With this data, I anticipate on determining and analyzing the aggregate data to learn more about the current trend of ride share services in New York City. While there are many routes to take when predicting future decisions with this data, I will mainly be focusing on predicting the tip amount based on the other factors that are involved with this table, such as the time, location, and distance of the trip.

Due to the fact that this data set contains so much information, even down to whether or not the passenger needed a wheelchair assisted vehicle or if the passenger agreed to having a shared car with other passengers, I anticipate learning more about the average behaviors of NYC TLC drivers and passengers. This data could also pair well with information regarding traffic patterns or vehicle collisions, or even any data that pertains to customer complaints on these rideshare apps.

The data is separated by month and year, and is stored as a parquet file, which can be used with pandas. The link to the data can be found below:

Uber/ Lyft (19 GB of data)-

[https://www.kaggle.com/datasets/jeffsinsel/nyc-fhvhv-data?select=fhvhv\\_tripdata\\_2020-06.parquet](https://www.kaggle.com/datasets/jeffsinsel/nyc-fhvhv-data?select=fhvhv_tripdata_2020-06.parquet)

<https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page>

## Milestone 2- Data Acquisition

Since I planned to download my data set from Kaggle, I installed the Kaggle CLI and input my username and key. The steps I used to download the files can be found in Appendix A. I repeated the code to download each of the individual files onto my EC2 instance and store the file onto my S3 “landing” bucket, replacing the bolded text with the appropriate file name. I also had to use curl to retrieve some of the more recent files directly from the NYC data repository.

My Amazon S3 landing bucket can be seen below:

The screenshot shows the Amazon S3 console interface for a bucket named 'landing/'. At the top right, there is a 'Copy S3 URI' button. Below the bucket name, there are two tabs: 'Objects' (selected) and 'Properties'. The main content area shows 'Objects (48)' with a description: 'Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)'. Below this, there are several action buttons: 'Refresh', 'Copy S3 URI', 'Copy URL', 'Download', 'Open', 'Delete', 'Actions', and 'Create folder'. There is also an 'Upload' button. A search bar with the placeholder 'Find objects by prefix' is present. Below the search bar is a table with columns: 'Name', 'Type', 'Last modified', 'Size', and 'Storage class'. The table lists three objects:

	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	<a href="#">fhvvhv_tripdata_2019-02.parquet</a>	parquet	September 27, 2023, 13:27:09 (UTC-04:00)	489.3 MB	Standard
<input type="checkbox"/>	<a href="#">fhvvhv_tripdata_2019-03.parquet</a>	parquet	September 27, 2023, 13:25:15 (UTC-04:00)	582.6 MB	Standard
<input type="checkbox"/>	<a href="#">fhvvhv_tripdata_2019-04.parquet</a>	parquet	September 27, 2023, 13:24:06 (UTC-04:00)	533.7 MB	Standard

### **Milestone 3- Exploratory Data Analysis**

The code for milestone 3 can be found in Appendix B, and the outputs of the code can be seen below:

```
#info on the fhv_df1 file

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 18479031 entries, 0 to 18479030

Data columns (total 24 columns):
#   Column                                Dtype
---  -
0   hvfhs_license_num                    object
1   dispatching_base_num                 object
2   originating_base_num                 object
3   request_datetime                     datetime64[ns]
4   on_scene_datetime                    datetime64[ns]
5   pickup_datetime                      datetime64[ns]
6   dropoff_datetime                     datetime64[ns]
7   PULocationID                         int64
8   DOLocationID                         int64
9   trip_miles                           float64
10  trip_time                             int64
11  base_passenger_fare                   float64
12  tolls                                 float64
13  bcf                                   float64
14  sales_tax                             float64
15  congestion_surcharge                  float64
16  airport_fee                           float64
17  tips                                  float64
```

```

18 driver_pay float64
19 shared_request_flag object
20 shared_match_flag object
21 access_a_ride_flag object
22 wav_request_flag object
23 wav_match_flag object

```

```
dtypes: datetime64[ns](4), float64(9), int64(3), object(8)
```

```
memory usage: 3.3+ GB
```

```
# describe the fhv_df1 file
```

	request_datetime	on_scene_datetime	pickup_datetime	dropoff_datetime	PULocationID	DOLocationID	trip_miles	trip_time	base_passenger_fare	toils	bcf
count	18479031	13587039	18479031	18479031	1.847903e+07	1.847903e+07	1.847903e+07	1.847903e+07	1.847903e+07	1.847903e+07	1.847903e+07
mean	2023-01-17 02:32:09.458061824	2023-01-17 06:15:32.572615936	2023-01-17 02:36:36.490860032	2023-01-17 02:54:51.087476736	1.393105e+02	1.426835e+02	4.870138e+00	1.094634e+03	2.155565e+01	1.034738e+00	6.819760e-01
min	2022-12-31 20:30:00	2022-12-31 21:23:03	2023-01-01 00:00:00	2023-01-01 00:02:27	1.000000e+00	1.000000e+00	0.000000e+00	0.000000e+00	-1.463400e+02	0.000000e+00	0.000000e+00
25%	2023-01-09 15:34:10	2023-01-09 19:35:55	2023-01-09 15:37:59	2023-01-09 15:58:02	7.500000e+01	7.600000e+01	1.550000e+00	5.750000e+02	1.043000e+01	0.000000e+00	3.100000e-01
50%	2023-01-17 09:35:41	2023-01-17 15:57:34	2023-01-17 09:39:34	2023-01-17 09:59:16	1.400000e+02	1.420000e+02	2.897000e+00	9.060000e+02	1.626000e+01	0.000000e+00	4.900000e-01
75%	2023-01-24 17:46:14	2023-01-24 20:15:07	2023-01-24 17:50:30	2023-01-24 18:10:02	2.110000e+02	2.200000e+02	6.015000e+00	1.404000e+03	2.617000e+01	0.000000e+00	8.100000e-01
max	2023-02-01 00:15:00	2023-01-31 23:59:53	2023-01-31 23:59:59	2023-02-01 01:47:23	2.650000e+02	2.650000e+02	4.075630e+02	3.535900e+04	1.455120e+03	1.843700e+02	6.471000e+01
std	NaN	NaN	NaN	NaN	7.510532e+01	7.803905e+01	5.655499e+00	7.453304e+02	1.801144e+01	3.757672e+00	6.096692e-01

```
# count of each variable in the fhv_df1 file
```

```

Out[4]: hvfhs_license_num      18479031

dispatching_base_num      18479031
originating_base_num      13587039
request_datetime          18479031
on_scene_datetime         13587039
pickup_datetime           18479031
dropoff_datetime          18479031
PULocationID              18479031
DOLocationID              18479031
trip_miles                 18479031
trip_time                  18479031
base_passenger_fare       18479031

```

```

tolls                18479031
bcf                  18479031
sales_tax            18479031
congestion_surcharge 18479031
airport_fee          18479031
tips                 18479031
driver_pay           18479031
shared_request_flag  18479031
shared_match_flag    18479031
access_a_ride_flag   18479031
wav_request_flag     18479031
wav_match_flag       18479031
dtype: int64

```

```
# list the variable names of the fhv_df1 file
```

```

['hvfhs_license_num', 'dispatching_base_num', 'originating_base_num',
'request_datetime', 'on_scene_datetime', 'pickup_datetime', 'dropoff_datetime',
'PULocationID', 'DOLocationID', 'trip_miles', 'trip_time',
'base_passenger_fare', 'tolls', 'bcf', 'sales_tax', 'congestion_surcharge',
'airport_fee', 'tips', 'driver_pay', 'shared_request_flag',
'shared_match_flag', 'access_a_ride_flag', 'wav_request_flag',
'wav_match_flag']

```

```
# count the number of rows with null values in the fhv_df2 file
```

```
>> Rows with null values: 4891992
```

---

```
#info on the fhv_df2 file
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20159102 entries, 0 to 20159101
```

```
Data columns (total 24 columns):
```

```

#      Column                Dtype

```

```

---  -----  -----
0   hvfhs_license_num    object
1   dispatching_base_num object
2   originating_base_num object
3   request_datetime     datetime64[ns]
4   on_scene_datetime    datetime64[ns]
5   pickup_datetime      datetime64[ns]
6   dropoff_datetime     datetime64[ns]
7   PULocationID         int64
8   DOLocationID         int64
9   trip_miles           float64
10  trip_time            int64
11  base_passenger_fare  float64
12  tolls               float64
13  bcf                 float64
14  sales_tax           float64
15  congestion_surcharge float64
16  airport_fee         object
17  tips               float64
18  driver_pay         float64
19  shared_request_flag object
20  shared_match_flag  object
21  access_a_ride_flag object
22  wav_request_flag   object
23  wav_match_flag     object

dtypes: datetime64[ns](4), float64(8), int64(3), object(9)

memory usage: 3.6+ GB

None

# describe the fhv_df2 file

```

	request_datetime	on_scene_datetime	pickup_datetime	dropoff_datetime	PULocationID	DOLocationID	trip_miles	trip_time	base_passenger_fare	tolls	bcf
count	20050204	13505053	20159102	20159102	2.015910e+07	2.015910e+07	2.015910e+07	2.015910e+07	2.015910e+07	2.015910e+07	2.015910e+07
mean	2019-02-15 02:09:02.991858432	2019-02-15 01:57:07.700399104	2019-02-15 00:36:08.897504256	2019-02-15 00:54:53.959078656	1.393497e+02	1.418818e+02	4.660525e+00	1.117965e+03	1.570783e+01	7.686299e-01	3.988125e-01
min	2019-01-31 23:19:44	2019-01-31 23:48:32	2019-02-01 00:00:00	2019-02-01 00:02:09	1.000000e+00	1.000000e+00	0.000000e+00	0.000000e+00	-1.632800e+02	0.000000e+00	0.000000e+00
25%	2019-02-08 08:28:40	2019-02-08 02:15:14	2019-02-08 06:07:44	2019-02-08 06:27:09	7.500000e+01	7.600000e+01	1.560000e+00	5.790000e+02	6.690000e+00	0.000000e+00	1.700000e-01
50%	2019-02-15 01:19:22.500000	2019-02-15 00:11:01	2019-02-14 23:29:07	2019-02-14 23:47:53	1.410000e+02	1.420000e+02	2.880000e+00	9.280000e+02	1.074000e+01	0.000000e+00	2.600000e-01
75%	2019-02-22 10:34:11	2019-02-22 15:12:38	2019-02-22 09:45:07	2019-02-22 10:04:03.750000128	2.110000e+02	2.190000e+02	5.670000e+00	1.453000e+03	1.898000e+01	0.000000e+00	4.800000e-01
max	2019-02-28 23:58:52	2019-02-28 23:59:50	2019-02-28 23:59:59	2019-03-01 06:02:46	2.650000e+02	2.650000e+02	4.692600e+02	8.384700e+04	1.097290e+03	1.710800e+02	2.793000e+01
std	NaN	NaN	NaN	NaN	7.521032e+01	7.743277e+01	5.415590e+00	7.730596e+02	1.612579e+01	3.185317e+00	4.517552e-01

# count of each variable in the fhv\_df2 file

```

Out[5]: hvfhs_license_num      20159102

dispatching_base_num      20158697

originating_base_num      14483914

request_datetime          20050204

on_scene_datetime         13505053

pickup_datetime           20159102

dropoff_datetime          20159102

PULocationID              20159102

DOLocationID              20159102

trip_miles                 20159102

trip_time                  20159102

base_passenger_fare        20159102

tolls                      20159102

bcf                        20159102

sales_tax                  20159102

congestion_surcharge      19646061

airport_fee                0

tips                       20159102

driver_pay                 20159102

shared_request_flag        20159102

shared_match_flag          20159102

```



```
access_a_ride_flag      20159102
wav_request_flag        20159102
wav_match_flag          0
dtype: int64
```

```
# list the variable names of the fhv_df2 file
```

```
['hvfhs_license_num', 'dispatching_base_num', 'originating_base_num',
 'request_datetime', 'on_scene_datetime', 'pickup_datetime', 'dropoff_datetime',
 'PULocationID', 'DOLocationID', 'trip_miles', 'trip_time',
 'base_passenger_fare', 'tolls', 'bcf', 'sales_tax', 'congestion_surcharge',
 'airport_fee', 'tips', 'driver_pay', 'shared_request_flag',
 'shared_match_flag', 'access_a_ride_flag', 'wav_request_flag',
 'wav_match_flag']
```

```
# count the number of rows with null values in the fhv_df2 file
```

```
Rows with null values: 20159102
```

```
# convert the dtype of the "trip_miles" variable to an integer, and ensure  
that this change was properly made
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20159102 entries, 0 to 20159101
```

```
Data columns (total 24 columns):
```

#	Column	Dtype
0	hvfhs_license_num	object
1	dispatching_base_num	object
2	originating_base_num	object
3	request_datetime	datetime64[ns]
4	on_scene_datetime	datetime64[ns]
5	pickup_datetime	datetime64[ns]
6	dropoff_datetime	datetime64[ns]
7	PULocationID	int64

```

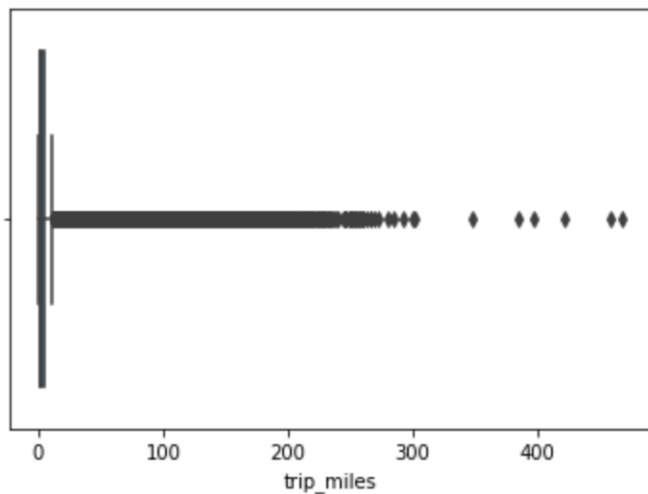
8   DOLocationID      int64
9   trip_miles         int64
10  trip_time          int64
11  base_passenger_fare float64
12  tolls              float64
13  bcf               float64
14  sales_tax          float64
15  congestion_surcharge float64
16  airport_fee        object
17  tips              float64
18  driver_pay         float64
19  shared_request_flag object
20  shared_match_flag  object
21  access_a_ride_flag object
22  wav_request_flag   object
23  wav_match_flag     object

```

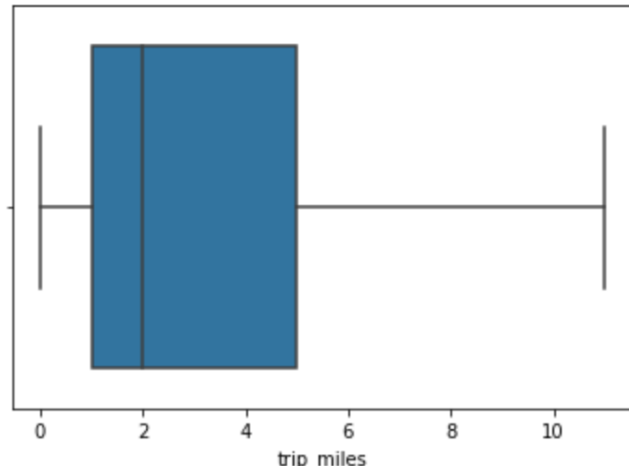
```
dtypes: datetime64[ns](4), float64(7), int64(4), object(9)
```

```
memory usage: 3.6+ GB
```

```
# create a boxplot showing the trip_miles for the fhv_df2 file
```



```
# create the same boxplot as above but remove the outliers for fhv_df2
```



From looking at the data, the sheer number of rows shows how many Uber and Lyft rides NYC uses in a given month- just looking at the February 2019 data and January 2023 data, the number of rides range from 18 million to 20 million. With this in mind, it did not surprise me to see how much variability can be in a single column. I chose to look at the “trip\_miles” column, which showcases the length of the Uber/ Lyft ride in miles. Only using a boxplot on the February 2019 data, it is hard to even see the data due to the amount of outliers that don’t fit the normal distribution of the column. When taking out the outliers, we can see that the normal distribution of trip miles ranges from 0-11 miles, which is miles away from the maximum of 469 miles. Due to the variability of the data, I’m expecting to deal with difficulties regarding visualizing the data in a way that it is easy for others to decipher. In addition, the amount of null values were concerning; the February 2019 data set had roughly 20 million rows with null values while the January 2023 data set had about more than 4 million rows with null values. The null values may be explained by the fact that much of the numerical data has been given the data type as float, so hopefully by converting the appropriate columns to integers would help to decrease the null values.

## Cleaning Data → Code in Appendix C

- The schema and head of my dataframe can be seen below:

```

hvfhs_license_num: string
dispatching_base_num: string
originating_base_num: string
request_datetime: timestamp
on_scene_datetime: timestamp
pickup_datetime: timestamp
dropoff_datetime: timestamp
PULocationID: long
DOLocationID: long
trip_miles: double
trip_time: double
base_passenger_fare: double
tolls: double
bcf: double
sales_tax: double
congestion_surcharge: double
airport_fee: double
tips: double
driver_pay: double
shared_request_flag: string
shared_match_flag: string
access_a_ride_flag: string

```

	hvh_license_num	dispatching_base_num	originating_base_num	request_datetime	on_scene_datetime	pickup_datetime	dropoff_datetime	PULocationID	DOLocationID	trip_miles	trip_time	base_passenger_fare	tolls	bcf	sales_tax	congestion_surcharge	airport_fee	tips	driver_pay	shared_request_flag	shared_match_flag	access_a_ride_flag	wav_request_flag	wav_match_flag
112	N	HV0003	B03404	B03404 2021-12-01 00:02:58	2021-12-01 00:05:38	2021-12-01 00:06:05	2021-12-01 00:19:05	80	N	2.39	13.0	13.27	0.0	0.4	1.18	0.0	0.0	0.0	9.41	N				
189	N	HV0003	B03404	B03404 2021-12-01 00:20:22	2021-12-01 00:21:26	2021-12-01 00:22:45	2021-12-01 00:43:47	112	N	4.91	[21.03333333333333]	22.59	0.0	0.68	2.0	0.0	0.0	5.05	16.23	N				
225	N	HV0003	B03404	B03404 2021-12-01 00:47:40	2021-12-01 00:48:50	2021-12-01 00:50:51	2021-12-01 00:59:20	49	N	1.59	8.483333333333333	9.58	0.0	0.29	0.85	0.0	0.0	0.0	6.26	N				
263	N	HV0003	B03404	B03404 2021-12-01 00:25:31	2021-12-01 00:29:07	2021-12-01 00:29:12	2021-12-01 00:37:22	239	N	1.78	8.166666666666666	9.15	0.0	0.21	0.63	2.75	0.0	7.0	6.07	N				

## Feature Engineering → Code in Appendix D

- The main goal of feature engineering was to make the data as digestible as possible before encoding it into a vector for my regression model, so there was much tweaking to be done. The specific changes I made are listed below:
  - Converting columns that had less than 5 standard values into flags (tolls, congestion\_charge, etc.)
  - Converting the taxi codes to the respective taxi companies
  - Converting the taxi zones to their respective boroughs
  - Drilling down request\_datetime into it's day of month, day type (weekend or weekday), month, year, and hour
  - Deriving how long the passenger waited for their taxi ride
  - Drop unnecessary columns
- The schema and head of the dataframe after feature engineering can be seen below:

```
jan23: pyspark.sql.dataframe.DataFrame
```

```
trip_miles: double
trip_time: double
base_passenger_fare: double
tips: double
driver_pay: double
shared_match_flag: string
congestion_surcharge_flag: string
tolls_flag: string
airport_fee_flag: string
hvfhs_company: string
PULocationBorough: string
DOLocationBorough: string
day_of_month: integer
day_type: string
month: integer
year: integer
request_hour: integer
wait_minutes: double
```

```
|trip_miles|trip_time|base_passenger_fare|tips|driver_pay|shared_match_flag|congestion_surcharge_flag|tolls_flag|airport_fee_flag|hvfhs_company|PULocationBorough|DOLocationBorough|day_of_month|day_type|month|year|request_hour|wait_minutes|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|0.94|28.483333333333334|25.95|15.22|27.83|N|Y|N|N|Uber|Manhattan|Manhattan|1|Weekend|1|2023|0|1.3|
```

attan	28.483333333333334	25.95	15.22	27.83	N	Y	N	N	Uber	Manhattan	Manhattan	1	Weekend	1	2023	0	1.3	
attan	34.483333333333334	60.14	0.0	50.15	N	Y	N	N	Uber	Manhattan	Manhattan	1	Weekend	1	2023	0		
ns	17.45	24.37	0.0	20.22	N	N	N	N	Uber	Manhattan	Queens	1	Weekend	1	2023	0		
ns	17.183333333333334	13.8	0.0	17.9	N	N	N	N	Uber	Manhattan	Queens	1	Weekend	1	2023	0		
ns	12.066666666666667	20.49	0.0	16.48	N	N	N	N	Uber	Manhattan	Queens	1	Weekend	1	2023	0		

```
|7.916666666666667|
```

## Modeling → Code in Appendix E

- To model my data, I combined all of the 2020 parquet files, as this was the year with the least FHV rides. Using an Indexer, Encoder, and Vector Assembler, I was able to create a pipeline that would predict whether or not the taxi driver had received a “good” tip, which I classified as 20% of the base\_passenger\_fare. After creating my model, I was able to test it on my testData (a random 30%) and validate this model across 3 folds.
- The AUC for a base tip of 20% or more was 0.692. This shows that the model was getting slightly better at predicting whether the driver would get a “good tip” as the threshold got higher, but the increase in AUC isn’t too substantial.
- My confusion matrix for a base tip of 20% can be seen below:

```
+-----+-----+-----+
|label|    0.0| 1.0|
+-----+-----+-----+
|  0.0|17601491| 306|
|  1.0| 832613|2301|
+-----+-----+-----+
accuracy:  0.94732520217838
precision:  0.9090909090909091
recall:    0.00010316191262186001
f1 score:  0.0002063004146638335
(0.94732520217838, 0.9090909090909091, 0.00010316191262186001, 0.0002063004146638335)
```

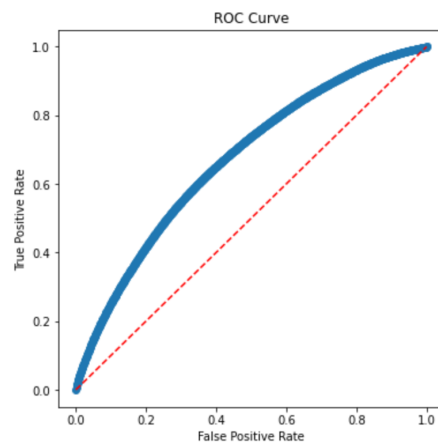
- The model was much better at determining true negatives (that someone left a “poor” tip) compared to true positives (leaving a “good” tip) for all tip thresholds. The F1 score is also very low across the thresholds, which is likely attributed to the low recall of the model.

## Milestone 5- Data Visualizing

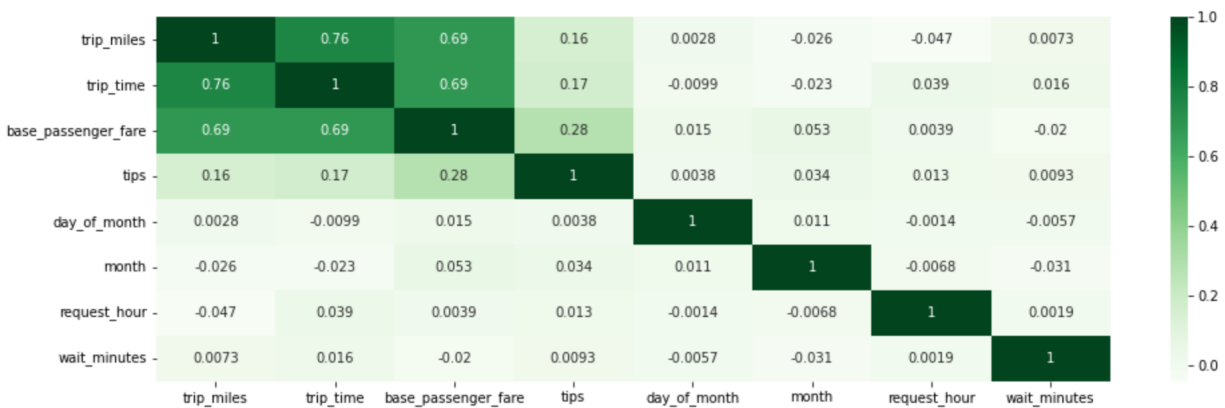
In an effort to visualize the accuracy of the models that I had created, I was able to develop ROC curves, which were able to tell me that the best model had an elasticNetParam of 0.0 and a regParam of 0.0 despite the different thresholds for tips. Through the ROC curves, we are able to see the tradeoff between accuracy and precision. The ROC curves are generally the same across the different thresholds. The code I used for data visualization can be found in Appendix F.

```
LogisticRegression_acc3ee1e7660__aggregationDepth 2
LogisticRegression_acc3ee1e7660__elasticNetParam 0.0
LogisticRegression_acc3ee1e7660__family auto
LogisticRegression_acc3ee1e7660__featuresCol features
LogisticRegression_acc3ee1e7660__fitIntercept True
LogisticRegression_acc3ee1e7660__labelCol label
LogisticRegression_acc3ee1e7660__maxBlockSizeInMB 0.0
LogisticRegression_acc3ee1e7660__maxIter 100
LogisticRegression_acc3ee1e7660__predictionCol prediction
LogisticRegression_acc3ee1e7660__probabilityCol probability
LogisticRegression_acc3ee1e7660__rawPredictionCol rawPrediction
LogisticRegression_acc3ee1e7660__regParam 0.0
LogisticRegression_acc3ee1e7660__standardization True
LogisticRegression_acc3ee1e7660__threshold 0.5
LogisticRegression_acc3ee1e7660__tol 1e-06
```

- The ROC curve for the 20% tip threshold can be seen below:

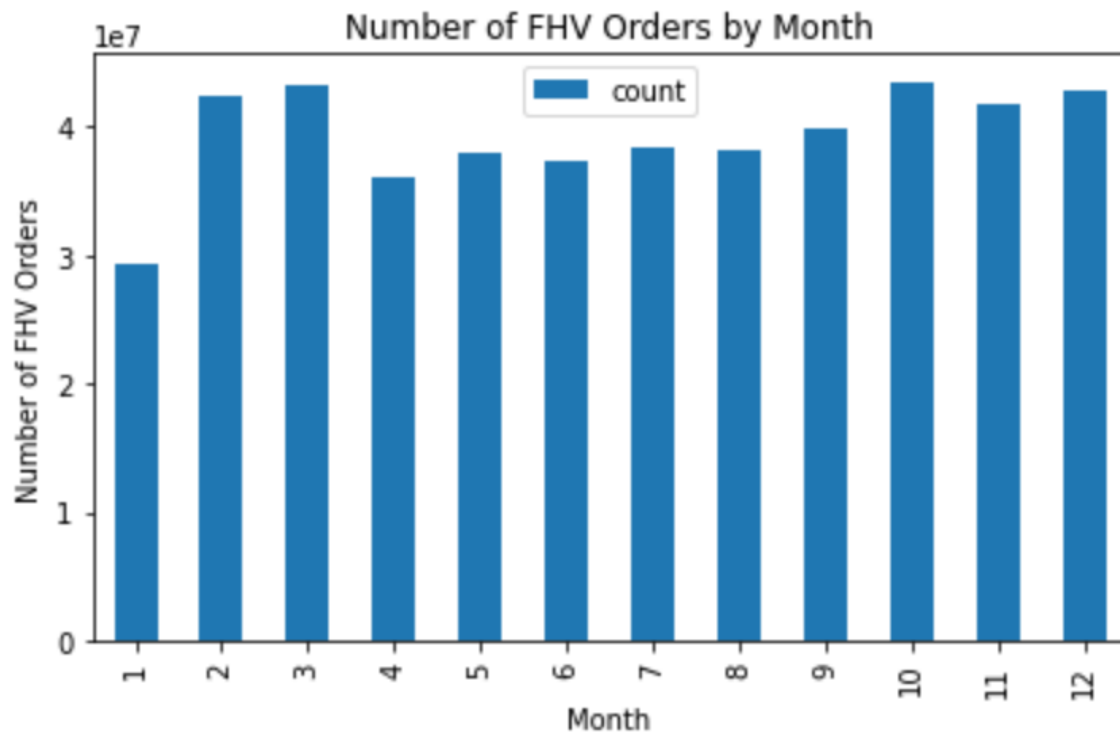
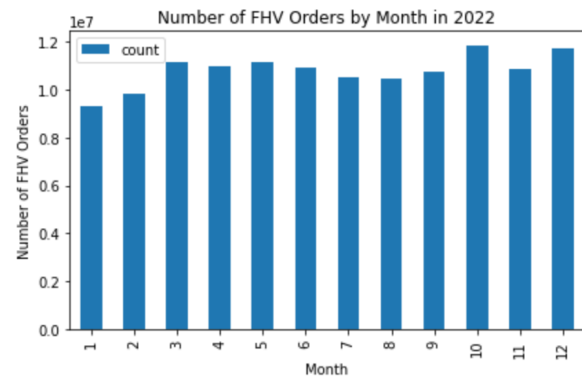
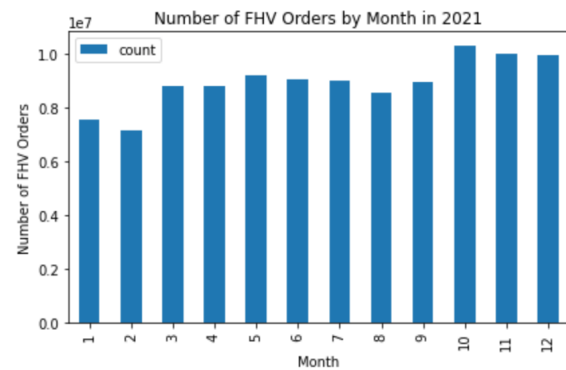
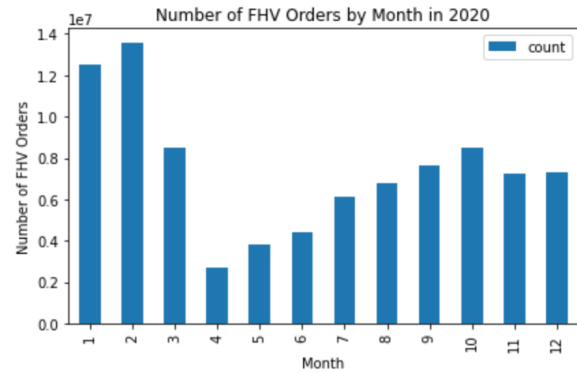
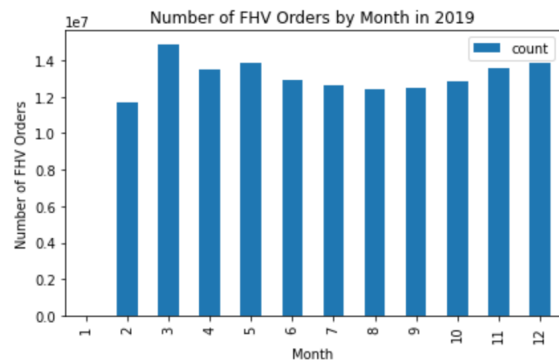


I had also developed a correlation matrix using all of the numeric values in the dataframe. From this correlation matrix, we are able to see that the day\_of\_month, month, request\_hour, and wait\_minutes have little to no correlation to the other variables. The factor that most affected “tips” would be the base\_passenger\_fare, though it is a very weak correlation. There are strong correlations between trip\_miles, trip\_time, and base\_passenger\_fare, which is to be expected as these variables are dependent on one another. The correlation matrix can be seen below:

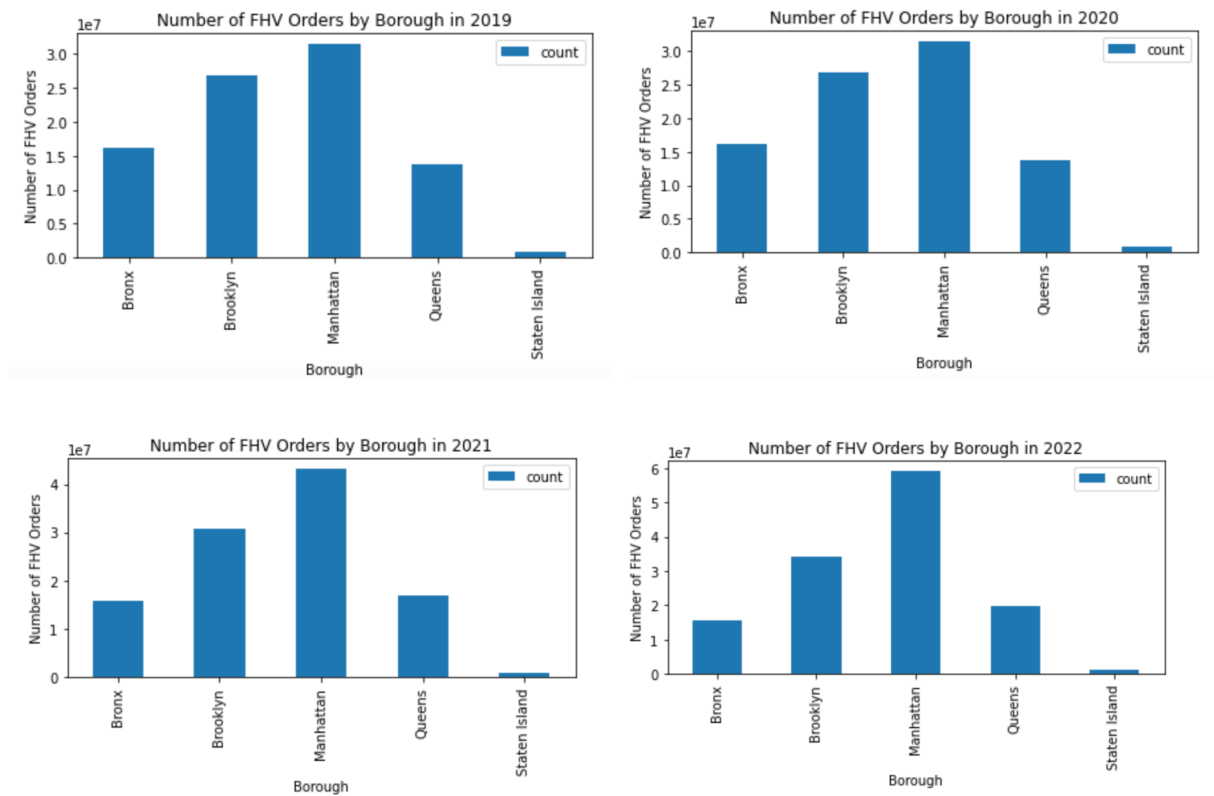


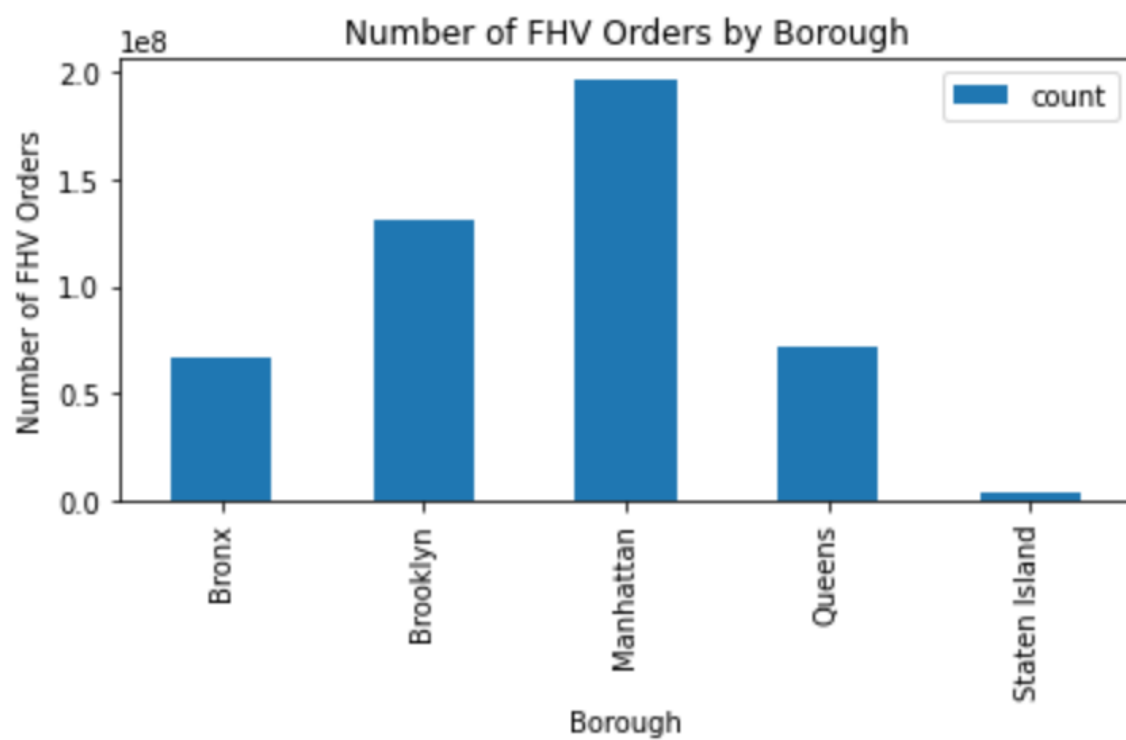
In addition, I was also able to visualize the number of rides that were requested by month, which I did per year as well as a cumulative total. We are able to see how FHV rides were greatly affected by the pandemic, as it was until October of 2021 that they were able to make pre-pandemic numbers again. By looking at the cumulative bar graph, we are able to see that February, March, October, and December would bring in the most rides from these FHV apps, whereas January and April bring in the lowest. The bar graphs can be seen below:





Lastly, I was able to develop bar graphs that showed the number of FHV rides per borough. Once again, I was able to show the bar graphs per year, as well as a cumulative view of the results. In all of the bar graphs, it is clear that Manhattan is the borough with the most FHV rides, which is to be expected due to the high population density and tourism that exists there. Brooklyn seems to be the next borough that has the highest number of FHV rides, followed by either the Bronx or Queens, tie for third place in the cumulative graph:





## **Milestone 6- Summary and Conclusions**

This machine learning project has helped me to understand the concept of big data, as well as learn the foundations of sources such as AWS, spark, and Databricks. A summary of my machine learning process can be seen below:

- Utilized Amazon EC2 to download and zip the parquet files from Kaggle, and load these files onto my S3 bucket.
- Showcased descriptive statistics using Databricks.
- Cleaned up my data in order to remove outliers and produce more meaningful columns.
- Performed feature engineering on Databricks, utilizing StringIndexer, OneHotEncoder, and VectorAssembler.
- Tested my data using a 70-30 split.
- Created visualizations to showcase my model performance, as well as visualizations to show the distribution of my data.

After completing my machine learning pipeline, I was able to draw conclusions about my model, as well as about the state of FHV rides. My model did not have the best overall performance, despite my best efforts at cleaning the data and removing outliers. However, I think this is due to the sheer amount of people who hadn't tipped, as this could result in heavily skewed data. This can be backed up by the fact that "tips" had a weak correlation to "base\_passenger\_fare." In addition, it didn't seem like any of the numerical variables had a strong correlation to the "tips" variable, which could also explain the poor model performance. If I were to redo this project, I would likely opt for a linear regression model instead of a logistic regression model and see if it changes the model performance.

<https://github.com/nazihamalik/fhv-tip-prediction/tree/main>

## **Appendix A- Code Used for Data Acquisition**

```
#look at the individual data files in the Kaggle data set
kaggle datasets files jeffsinsel/nyc-fhvhv-data

#download individual data files onto EC2
kaggle datasets download -d jeffsinsel/nyc-fhvhv-data -f
fhvhv_tripdata_2019-06.parquet

#unzip file:
unzip fhvhv_tripdata_2019-06.parquet

#copy file onto Amazon S3 Bucket named "landing":
aws s3 cp fhvhv_tripdata_2019-06.parquet
s3://my-project-nm/landing/fhvhv_tripdata_2019-06.parquet

#remove the downloaded file off of EC2 once confirming it was uploaded onto my
bucket:

rm fhvhv_tripdata_2019-06.parquet
rm fhvhv_tripdata_2019-06.parquet.zip


#download individual data files onto EC2
kaggle datasets download -d jeffsinsel/nyc-fhvhv-data -f
fhvhv_tripdata_2019-02.parquet

#unzip file:
unzip fhvhv_tripdata_2019-02.parquet

#copy file onto Amazon S3 Bucket named "landing":
aws s3 cp fhvhv_tripdata_2019-02.parquet
s3://my-project-nm/landing/fhvhv_tripdata_2019-02.parquet

#remove the downloaded file off of EC2 once confirming it was uploaded onto my
bucket:

rm fhvhv_tripdata_2019-02.parquet
```

```
rm fhvhv_tripdata_2019-02.parquet.zip
```

```
#download individual data files onto EC2
```

```
kaggle datasets download -d jeffsinsel/nyc-fhvhv-data -f
```

```
fhvhv_tripdata_2019-03.parquet
```

```
#unzip file:
```

```
unzip fhvhv_tripdata_2019-03.parquet
```

```
#copy file onto Amazon S3 Bucket named "landing":
```

```
aws s3 cp fhvhv_tripdata_2019-03.parquet
```

```
s3://my-project-nm/landing/fhvhv_tripdata_2019-03.parquet
```

```
#remove the downloaded file off of EC2 once confirming it was uploaded onto my  
bucket:
```

```
rm fhvhv_tripdata_2019-03.parquet
```

```
rm fhvhv_tripdata_2019-03.parquet.zip
```

```
#download the data set using curl, and with the pipe, I was able to upload it  
onto my S3 bucket:
```

```
curl -SL
```

```
https://d37ci6vzurychx.cloudfront.net/trip-data/fhvhv\_tripdata\_2022-01.parquet
```

```
| aws s3 cp - s3://my-project-nm/landing/fhvhv_tripdata_2022-01.parquet
```

```
#download the data set using curl, and with the pipe, I was able to upload it  
onto my S3 bucket:
```

```
curl -SL
```

```
https://d37ci6vzurychx.cloudfront.net/trip-data/fhvhv\_tripdata\_2020-12.parquet
```

```
| aws s3 cp - s3://my-project-nm/landing/fhvhv_tripdata_2020-12.parquet
```

## **Appendix B- Code Used for Exploratory Data Analysis**

```
!pip install fsspec s3fs boto3
!pip install pyarrow fastparquet
!pip install seaborn

pip install --upgrade pandas

import pandas as pd
import pyarrow
import fastparquet
import seaborn as sns

fhv_df1 =
pd.read_parquet("https://my-project-nm.s3.us-east-2.amazonaws.com/landing/fhvhv
_tripdata_2023-01.parquet")

# column names and data types for each column
print(fhv_df1.info())

# summary statistics of each column
fhv_df1.describe()

# number of observations
fhv_df1.count()

# list of variable/ column names
variable_list = list(fhv_df1)
print(variable_list)

# null values?
# Which columns have nulls or NaN ?
fhv_df1.columns[fhv_df1.isnull().any()].tolist()
# How many rows have nulls?
print("Rows with null values:", fhv_df1.isnull().any(axis=1).sum())
```

```

fhv_df2 =
pd.read_parquet("https://my-project-nm.s3.us-east-2.amazonaws.com/landing/fhvhv\_tripdata\_2019-02.parquet")

# column names and data types for each column
print(fhv_df2.info())

# summary statistics of each column
fhv_df2.describe()

# number of observations
fhv_df2.count()

# list of variable/ column names
variable_list = list(fhv_df2)
print(variable_list)

# null values?
# Which columns have nulls or NaN ?
fhv_df2.columns[fhv_df2.isnull().any()].tolist()
# How many rows have nulls?
print("Rows with null values:", fhv_df2.isnull().any(axis=1).sum())

# convert variable trip_miles from float to integer
fhv_df2['trip_miles'] = fhv_df2['trip_miles'].fillna(0).astype(int)

# make sure that the changed dtype is stored
print(fhv_df2.info())

# make a boxplot showing the trip_miles
sns.boxplot(x=fhv_df2["trip_miles"])

# make a boxplot without outliers for trip_miles
sns.boxplot(x=fhv_df2["trip_miles"], showfliers=False)

```



## **Appendix C- Code Used for Cleaning and Normalizing the Data**

```
# To work with Amazon S3 install boto3, s3fs
%pip install "boto3>=1.28" "s3fs>=2023.3.0"
# If your files are in Parquet format, install pyarrow and fastparquet
%pip install pyarrow fastparquet
# For visualizations, install seaborn
%pip install seaborn

Spark

import os

# To work with Amazon S3 storage, set the following variables using your AWS
Access Key and Secret Key
# Set the Region to where your files are stored in S3.
access_key = '#####'
secret_key = '#####'
# Set the environment variables so boto3 can pick them up later
os.environ['AWS_ACCESS_KEY_ID'] = access_key
os.environ['AWS_SECRET_ACCESS_KEY'] = secret_key encoded_secret_key =
secret_key.replace("/", "%2F")
aws_region = "us-east-2"

sc._jsc.hadoopConfiguration().set("fs.s3a.access.key", access_key)
sc._jsc.hadoopConfiguration().set("fs.s3a.secret.key", secret_key)
sc._jsc.hadoopConfiguration().set("fs.s3a.endpoint", "s3." + aws_region +
".amazonaws.com")

from pyspark.sql import SparkSession
from pyspark.sql import functions as F
from pyspark.sql.types import DoubleType
from pyspark.sql.functions import col, lit, when, date_format, expr, to_date,
to_timestamp, date_format, split, hour, dayofweek, dayofmonth, month, year

### DATA CLEANING
# read the data from my landing folder
dec21 =
spark.read.parquet("s3://my-project-nm/landing/fhvhv_tripdata_2021-12.parquet")
```

```
# drop necessary null values
columns_to_drop_null = ["request_datetime", "on_scene_datetime",
"pickup_datetime", "dropoff_datetime" ]
dec21 = dec21.na.drop(subset=columns_to_drop_null)

# convert trip_time from seconds to minutes
dec21 = dec21.withColumn('trip_time', (F.col('trip_time') /60))

# drop records where trip_miles is an outlier based on previous EDA
trip_miles_condition = col("trip_miles") <= 15
dec21 = dec21.where(trip_miles_condition)

dec21.show(5)

# write the parquet to my S3 Raw Bucket
dec21.write.parquet("s3://my-project-nm/raw/cleaned_fhvhv_tripdata_2021-12.parquet")
```

## **Appendix D- Code Used for Feature Engineering**

```
###FEATURE ENGINEERING

jan23 =
spark.read.parquet("s3://my-project-nm/raw/cleaned_fhv_hv_tripdata_2023-01.parquet")

### convert congestion_surcharge to a flag
jan23 = jan23.withColumn("congestion_surcharge_flag",
when(jan23.congestion_surcharge >0, "Y")
.otherwise("N"))

)

### convert tolls to a flag
jan23 = jan23.withColumn("tolls_flag",
when(jan23.tolls >0, "Y")
.otherwise("N"))

)

### convert airport_fee to a flag
jan23 = jan23.withColumn("airport_fee_flag",
when(jan23.airport_fee>0, "Y")
.otherwise("N"))

)

### convert hvfhs_license_num to the respected fhv company
jan23 = jan23.withColumn("hvfhs_company",
when(jan23.hvfhs_license_num== "HV0005", "Lyft")
.when(jan23.hvfhs_license_num== "HV0002", "Juno")
.when(jan23.hvfhs_license_num== "HV0003", "Uber")
.when(jan23.hvfhs_license_num == "HV0004", "Via")
.otherwise("N/A"))

)

### convert PULocationID to the respected borough
jan23 = jan23.withColumn("PULocationBorough",
when(jan23.PULocationID.isin(Bronx), "Bronx")
.when(jan23.PULocationID.isin(Brooklyn), "Brooklyn")
```

```

.when(jan23.PULocationID.isin(Queens), "Queens")
.when(jan23.PULocationID.isin(Manhattan), "Manhattan")
.when(jan23.PULocationID.isin(Staten_Island), "Staten Island")
.otherwise("N/A")
)

# drop records if the PULocationBorough isn't registered
PULocationCondition = col("PULocationBorough") != "N/A"
jan23 = jan23.where(PULocationCondition)

### convert DOLocation ID to the respected borough
jan23 = jan23.withColumn("DOLocationBorough",
when(jan23.DOLocationID.isin(Bronx), "Bronx")
.when(jan23.DOLocationID.isin(Brooklyn), "Brooklyn")
.when(jan23.DOLocationID.isin(Queens), "Queens")
.when(jan23.DOLocationID.isin(Manhattan), "Manhattan")
.when(jan23.DOLocationID.isin(Staten_Island), "Staten Island")
.otherwise("N/A")
)

# drop records if the DOLocationBorough isn't registered
DOLocationCondition = col("DOLocationBorough") != "N/A"
jan23 = jan23.where(DOLocationCondition)

### separate request_datetime into date and time columns
jan23 = jan23.withColumn("request_date", to_date("request_datetime",
"yyyy-MM-dd"))
jan23 = jan23.withColumn("day_of_month", dayofmonth(col("request_date")))
jan23 = jan23.withColumn("day_of_week", dayofweek(col("request_date")))
def is_weekday(day):
    return "Weekday" if day in range(2, 7) else "Weekend"
spark.udf.register("is_weekday", is_weekday)
jan23 = jan23.withColumn("day_type",
col("day_of_week").cast("int").alias("day_type"))
jan23 = jan23.withColumn("day_type", expr("is_weekday(day_type)"))

jan23 = jan23.withColumn("month", month(col("request_datetime")))
jan23 = jan23.withColumn("year", year(col("request_datetime")))

split_col = split(jan23['request_datetime'], ' ')

```

```

jan23 = jan23.withColumn('request_time', split_col.getItem(1))
jan23 = jan23.withColumn("request_hour", hour(col("request_time")))

## create a new column to show the wait time of a passenger
jan23 = jan23.withColumn("wait_seconds",
(col("on_scene_datetime").cast("long")- col("request_datetime").cast("long")))
jan23 = jan23.withColumn("wait_minutes", col("wait_seconds") / 60)

# drop unnecessary columns
columns_to_drop= ["access_a_ride_flag", "wav_request_flag", "wav_match_flag",
"bcf", "sales_tax", "shared_request_flag", "dispatching_base_num",
"originating_base_num", "DOLocationID", "PULocationID", "hvfhs_license_num",
"airport_fee", "tolls", "congestion_surcharge", "request_datetime",
"on_scene_datetime", "pickup_datetime", "dropoff_datetime", "wait_seconds",
"request_time", "request_date", "day_of_week"]
jan23 = jan23.drop(*columns_to_drop)

jan23.show(5, truncate=False)

# write the parquet to my S3 Trusted Bucket
jan23.write.parquet("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2023-01.
parquet")

```

## **Appendix E- Code Used for Modeling Pipeline**

```
# read 2020 files
parquet_2020 =
[("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-01.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-02.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-03.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-04.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-05.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-06.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-07.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-08.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-09.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-10.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-11.parquet"),
("s3://my-project-nm/trusted/trusted_fhvhv_tripdata_2020-12.parquet")]

fhv2020 = [spark.read.parquet(p) for p in parquet_2020]
combined_2020 = fhv2020[0]
for p in fhv2020[1:]:
    combined_2020 = combined_2020.union(p)

# create new columns to show our tip label
combined_2020 = combined_2020.withColumn("label20",
(combined_2020.base_passenger_fare*0.20))

# COMPARING TIPS THAT ARE 20%
# Create a label. =1 if good tip, =0 otherwise
combined_2020 = combined_2020.withColumn("label", when(combined_2020.tips >=
combined_2020.label20, 1.0).otherwise(0.0))

# Create an indexer for the string based columns
indexer = StringIndexer(inputCols=["shared_match_flag",
"congestion_surcharge_flag", "tolls_flag", "airport_fee_flag", "hvfhs_company",
"PULocationBorough", "DOLocationBorough", "day_type"],
    outputCols=["shared_matchIndex", "congestionIndex",
"tollsIndex", "airportIndex", "hvfhsIndex", "PUIndex", "DOIndex",
"daytypeIndex"])
```

```

# Create an encoder for the indexes and the integer columns.
encoder = OneHotEncoder(inputCols=["shared_matchIndex", "congestionIndex",
"tollsIndex", "airportIndex", "hvfhsIndex", "PUIndex", "DOIndex",
"daytypeIndex", "day_of_month", "month", "year", "request_hour"],
                        outputCols=["shared_matchVector", "congestionVector",
"tollsVector", "airportVector", "hvfhsVector", "PUVector", "DOVector",
"daytypeVector", "daymonthVector", "monthVector", "yearVector",
"requesthourVector"], dropLast=True, handleInvalid="keep")

# Create an assembler for the individual feature vectors and the float/double
columns
assembler = VectorAssembler(inputCols=["shared_matchVector",
"congestionVector", "tollsVector", "airportVector", "hvfhsVector", "PUVector",
"DOVector", "daytypeVector", "daymonthVector", "monthVector", "yearVector",
"requesthourVector", "trip_miles", "trip_time", "base_passenger_fare",
"wait_minutes"], outputCol="features")

# Create a LogisticRegression Estimator
lr = LogisticRegression()

# split the data into two subsets
trainingData, testData = combined_2020.randomSplit([0.7, 0.3])
# create the pipeline
fhv_pipe = Pipeline(stages=[indexer, encoder, assembler, lr])

# Create the parameter grid
grid = ParamGridBuilder()
grid = grid.addGrid(lr.regParam, [0.0, 0.5, 1.0])
grid = grid.addGrid(lr.elasticNetParam, [0, 1])
grid = grid.build()

print("number of models to be tested: ", len(grid))
# Create a BinaryClassificationEvaluator to evaluate how well the model works
evaluator = BinaryClassificationEvaluator(metricName="areaUnderROC")
# Create the CrossValidator using the hyperparameter grid
cv = CrossValidator(estimator=fhv_pipe,
                    estimatorParamMaps=grid,

```

```
        evaluator=evaluator,  
        numFolds=3)  
  
# Train the models  
cv = cv.fit(trainingData)  
predictions = cv.transform(testData)  
auc = evaluator.evaluate(predictions)  
print("auc:", auc)
```



## **Appendix F- Code Used for Data Visualization**

```
# Show the confusion matrix
predictions.groupby('label').pivot('prediction').count().sort('label').show()

# Save the confusion matrix
cm =
predictions.groupby('label').pivot('prediction').count().fillna(0).collect()
def calculate_recall_precision(cm):
    tn = cm[0][1] #true negative
    fp = cm[0][2] #false positive
    fn = cm[1][1] #false negative
    tp = cm[1][2] #true positive
    precision = tp / ( tp + fp )
    recall = tp / ( tp + fn )
    accuracy = ( tp + tn ) / ( tp + tn + fp + fn )
    f1_score = 2 * ( ( precision * recall ) / ( precision + recall ) )
    print("accuracy: ", accuracy)
    print("precision: ", precision)
    print("recall: ", recall)
    print("f1 score: ", f1_score)
    return accuracy, precision, recall, f1_score

print( calculate_recall_precision(cm) )

## SHOW ROC CURVES
# Look at the parameters for the best model that was evaluated from the grid
parammap = cv.bestModel.stages[3].extractParamMap()

for p, v in parammap.items():
    print(p, v)

# Grab the model from Stage 3 of the pipeline
mymodel = cv.bestModel.stages[3]

import matplotlib.pyplot as plt
plt.figure(figsize=(5,5))
plt.plot(mymodel.summary.roc.select('FPR').collect(),
         mymodel.summary.roc.select('TPR').collect())
```

```

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title("ROC Curve")
plt.savefig("roc1.png")

parammap = cv.bestModel.stages[3].extractParamMap()

for p, v in parammap.items():
    print(p, v)

# Grab the model from Stage 3 of the pipeline
mymodel = cv.bestModel.stages[3]
plt.figure(figsize=(6,6))
plt.plot([0, 1], [0, 1], 'r--')
x = mymodel.summary.roc.select('FPR').collect()
y = mymodel.summary.roc.select('TPR').collect()
plt.scatter(x, y)
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title("ROC Curve")
plt.savefig("reviews_roc.png")

## CORRELATION MATRIX
correlation_columns = ['trip_miles', 'trip_time', 'base_passenger_fare',
'tips', 'day_of_month', 'month', 'request_hour', 'wait_minutes']
numeric_all = spark.read.parquet(*parquet_all).select(correlation_columns)

# Convert the numeric values to vector columns
vector_column = "correlation_features"

# Make a list of all of the numeric columns
numeric_columns = ['trip_miles', 'trip_time', 'base_passenger_fare', 'tips',
'day_of_month', 'month', 'request_hour', 'wait_minutes']

# Use a vector assembler to combine all of the numeric columns together
assembler = VectorAssembler(inputCols=numeric_columns, outputCol=vector_column)
sdf_vector = assembler.transform(numeric_all).select(vector_column)

```

```

# Create the correlation matrix, then get just the values and convert to a list
matrix = Correlation.corr(sdf_vector, vector_column).collect()[0][0]
correlation_matrix = matrix.toArray().tolist()

# Convert the correlation to a Pandas dataframe
correlation_matrix_df = pd.DataFrame(data=correlation_matrix,
columns=numeric_columns, index=numeric_columns)

# Create the plot using Seaborn
plt.figure(figsize=(16,5))
sns.heatmap(correlation_matrix_df,
            xticklabels=correlation_matrix_df.columns.values,
            yticklabels=correlation_matrix_df.columns.values,
            cmap="Greens",
            annot=True)
plt.savefig("correlation_matrix.png")

## BAR GRAPHS OF NUMBER OF FHV ORDERS BY MONTH
# select the columns that are needed
month_columns = ["month", "request_hour", "tips"]

# create a sdf with the specified parquet files and columns
monthrequesthourtips = spark.read.parquet(*parquet_all).select(month_columns)

# Use groupby to get a count by date. Then convert to pandas dataframe
month = monthrequesthourtips.groupby("month").count().sort("month").toPandas()

# Using Pandas built-in plotting functions
# Create a bar plot using the columns order_date and count
monthplot = month.plot.bar('month', 'count')
# Set the x-axis and y-axis labels
monthplot.set(xlabel='Month', ylabel='Number of FHV Orders')
# Set the title
monthplot.set(title='Number of FHV Orders by Month')
monthplot.figure.set_tight_layout('tight')
# Save the plot as a PNG file
monthplot.get_figure().savefig("order_rides_by_month.png")

```

```
## BAR GRAPH OF NUMBER OF FHV RIDES BY BOROUGH
# select the columns that are needed
borough_columns = ["PULocationBorough", "DOLocationBorough", "tips"]

# create a sdf with the specified parquet files and columns
boroughtips = spark.read.parquet(*parquet_all).select(borough_columns)

# Use groupby to get a count by date. Then convert to pandas dataframe
borough =
boroughtips.groupby("PULocationBorough").count().sort("PULocationBorough").toPa
ndas()

# Using Pandas built-in plotting functions
# Create a bar plot using the columns order_date and count
boroughplot = borough.plot.bar('PULocationBorough', 'count')
# Set the x-axis and y-axis labels
boroughplot.set(xlabel='Borough', ylabel='Number of FHV Orders')
# Set the title
boroughplot.set(title='Number of FHV Orders by Borough')
boroughplot.figure.set_tight_layout('tight')
# Save the plot as a PNG file
boroughplot.get_figure().savefig("order_rides_by_borough.png")
```